

PG Software Lab - CSP 509
Major-project Specification Phase A
Due date: Oct 23, 2019 11:55pm
Total Weightage 23%

Instructions:

- All submissions must be made through the Moodle site for this course
- Only one submission per team would be considered and graded. It would be assumed that all members of the team have participated equally and same score would be given to all members of the team.
- **Your submission should have names of all the members of your team.**
- Any assumptions made while solving the problem should be clearly stated in the solution.
- **You can take-up this project in teams of size 2 or size 3. Teams size-3 would have to do more work in Phase B of this project. You are not allowed to different teams across phase A and phase B of the project.**
- As always correctness of the algorithm must be ensured.
- TAs would be quizzing you on your code. You must understand each and every line of your submitted code. Also the implementation specifications mentioned in the questions need to be strictly followed. Failure to adhere to these requirements would result in substantial loss of points.
- Also pay attention to the scalability of the code. Your code would be tested on large datasets.
- **Very Important: Your code should not have a directory structure. All files (code + dataset + written material for questions) should be present in just one folder. Note that this is absolutely crucial for grading this assignment.**

In this phase, you are required to implement Extendible Hashing. Specialized libraries for managing hash tables must not be used. You may use existing libraries for basic data structures like vectors and associative array (e.g, map in C++ STL libraries or equivalent in JAVA). You may also use math libraries as needed for tasks like, generating random numbers, converting to and from binary format, etc.

Dataset creation:

For this question, you would have to create a synthetic table (simulating sales records of department stores) containing 1 Lakh records. Each record in this file contains four fields: (1) Transaction ID (an integer), (2) Transaction sale amount (an integer), (3) Customer name (string) and, (4) category of item. Transaction ID is an integer to identify each transaction uniquely in the dataset. Transaction sale amount is a random integer between 1 and 500000. Customer name is a random 3 letter string. You can model this as a character array of length 3. Category of the item is a random integer between 1 --1500 (uniformly distributed across the table).

After creating this dataset, you need to simulate its storage on a disk. Define a disk block as a file which can store only 300 records of the synthetic table. Assume unspanned organization i.e., records are not allowed to span across two disk block. Following this store your entire synthetic table as collection of these “disk blocks.” Each disk block (simulated as a file) should have a unique name, for that you can name them as 1, 2, 3, ..., etc. Basically, your original synthetic sales table would be stored as a series of files. The first disk block (file) would store Row 1 -- Row 300; second disk block (second file) would store Row 301 – Row 600. For sake of simplicity use text files for simulating the disk blocks. Also note that each disk block should have an entry at its end which stores the file name of the next disk block for the file. Additionally, in your code you should store the file name of the first disk block of the synthetic sales table. This would be needed a later stage to perform a linear scan through the file.

After creating the dataset, you are required to index the file using an Extendible hash. Basically, in your extendible hash you would be storing the following kinds of records: {TransactionID, <Filename of the disk block which stores this record>}. Implementation specification of the extendible hash is given below:

Details of Extendible Hash:

- We would be storing index records of the format “{TransactionID, <Filename of the disk block which stores this record>}” in the extendible hash.
- We would hashing on “TransactionID.”
- Ideally, the buckets in the extendible hash should be stored in the secondary memory. However, for the purpose of this project, they would be stored in something called “Simulated Secondary Memory (detailed below).”
- The directory or bucket address table of the extendible hash would contain hash prefix and pointer to the bucket sitting in “Simulated Secondary Memory (detailed below).”

Simulated Secondary Memory:

Following tips would help you achieve this.

- (a) The secondary memory can be simulated through an array of the abstract data-type “bucket”. You can fix a very large size for this array.
- (b) The bucket capacity is fixed in terms of number of index records it can contain. Do not hard code this number as it would be varied in the experiments.
- (c) Indices in this array form our “bucket address / hardware address.”
- (d) Here, the bucket abstract data-type would have the following information:
 - a. Number of empty spaces
 - b. An array of structures to store the index records. Length of this array is fixed according to the parameter “bucket-size” specified.
 - c. Link to the next bucket (valid only if this bucket is overflowing)
 - d. All buckets in the overflow chain must be linked. The last bucket of the overflow chain must have a special character denoting that it is end of the overflow chain.
- (e) **Note that in your entire code, there should be only one abstract data type for bucket for all the purposes, viz., records, overflow buckets for both records and directory entries.**
- (f) It is advisable to keep a separate area in the secondary memory for storing the overflow buckets.
- (g) **There should be only one instance of secondary memory running in your code.**

Main Memory

- (a) You can assume to have enough “main memory” for “bringing” in the buckets to be rehashed.
- (b) In addition to item (a) “Main memory” can hold upto 1024 directory entries. The rest resides in “Simulated Secondary memory.”
- (c) **Directory entries overflowing into secondary memory would be using the same bucket abstract data type which was previously declared in item “Simulated Secondary Memory”.**

Other Details:

- (a) The most significant bits are extracted to find the directory entry.
- (b) Only one directory expansion is allowed per record insertion. Following the directory expansion, you may attempt to resolve the collision (if it still persists) by increasing the local depth (if local depth < global depth). In case the collision is still not resolved, just create an overflow bucket.
- (c) “Main memory” can hold upto 1024 directory entries. The rest resides in “Simulated Secondary Memory.”

Evaluation Details:

- (1) TAs would use a small dataset to check the running of your code. For this, they may give you a file containing 15-20 records. Your code should read this file and, break it into disk blocks (simulated as files) and then load the index records in the extendible hash.
- (2) Extendible hash must have a separate insert function which would insert any given arbitrary “index record” into the extendible hash.
- (3) Extendible hash must have a separate intuitive visualize function for checking the correctness of the code.